

Physics-Informed Neural Networks for Aerodynamic Analysis of Rotating Cylinder Flows

*Thesis is submitted in the partial fulfillment
of the requirements for the degree of*

BACHELOR OF TECHNOLOGY

In

MECHANICAL ENGINEERING

By

SHUBHAM KUMAR (22ME8107)

RAHUL YADAV (22ME8064)

Under the guidance

of

Dr. SAIF AKRAM

Assistant Professor

Department of Mechanical Engineering



National Institute of Technology Durgapur

West Bengal, India

MAY 2026



**DEPARTMENT OF MECHANICAL ENGINEERING
NATIONAL INSTITUTE OF TECHNOLOGY
DURGAPUR, INDIA**

DECLARATION

We the undersigned declare that the thesis work entitled “**Physics-Informed Neural Networks for Aerodynamic Analysis of Rotating Cylinder Flows**”, submitted towards partial fulfillment of requirements for the award of the degree in **Bachelor of Technology in Mechanical Engineering** is our original work and this declaration does not form the basis for award of any degree or any similar title to the best of our knowledge.

Durgapur
May, 2026

Shubham Kumar (22ME8107)
Rahul Yadav (22ME8064)



**DEPARTMENT OF MECHANICAL ENGINEERING
NATIONAL INSTITUTE OF TECHNOLOGY
DURGAPUR, INDIA**

CERTIFICATE OF RECOMMENDATION

This is to certify that the thesis entitled “**Physics-Informed Neural Networks for Aerodynamic Analysis of Rotating Cylinder Flows**”, submitted by **Shubham Kumar & Rahul Yadav** of Department of Mechanical Engineering, National Institute of Technology, Durgapur, in partial fulfillment of the requirements for the award of the degree in **Bachelor of Technology in Mechanical Engineering** is a bonafide record of work carried out by them under my guidance during the academic year 2025 – 2026.

Dr. Asit Baran Puri
Head of the Department
Department of Mechanical Engineering
National Institute of Technology
Durgapur

Dr. Saif Akram
Assistant Professor
(Project Supervisor)
Department of Mechanical Engineering
National Institute of Technology
Durgapur



DEPARTMENT OF MECHANICAL ENGINEERING
NATIONAL INSTITUTE OF TECHNOLOGY
DURGAPUR, INDIA

CERTIFICATE OF APPROVAL

This is to certify that we have examined the thesis entitled “**Physics-Informed Neural Networks for Aerodynamic Analysis of Rotating Cylinder Flows**”, submitted by **Shubham Kumar & Rahul Yadav** and hereby accord our approval of it as a study carried out and presented in a manner required for its acceptance in partial fulfillment of the requirements for the award of the degree in **Bachelor of Technology in Mechanical Engineering** for which it has been submitted. It is to be understood that by this approval the undersigned do not necessarily endorse or approve any statement made, opinion expressed or conclusion drawn therein but approve the thesis only for the purpose for which it is submitted.

Examiners:

Name

Signature

ACKNOWLEDGEMENTS

I express my sincere gratitude to the National Institute of Technology Durgapur for providing the academic environment, computational resources, and structured undergraduate programme within which this thesis was conceived and completed. The Institute's emphasis on original research and its integration of advanced computational methods into the curriculum made a project of this scope achievable at the B.Tech. level.

I am deeply thankful to my thesis supervisor, Dr. Saif Akram, for patient guidance, technical depth, and regular feedback that shaped this work from its initial concept to its final form. The direction to ground the machine-learning approach in specific, well-validated CFD literature — rather than abstract benchmark datasets — proved fundamental to the quality and physicality of the trained model.

I acknowledge with gratitude the original researchers whose published numerical data provided the empirical foundation of this thesis: Dr. S.J. Karabelas (NTUA), whose 2010 LES study at $Re = 140,000$ provided the primary validation dataset; and Karabelas, Koumroglou, Argyropoulos, and Markatos, whose 2012 RANS study extended the training domain to supercritical Reynolds numbers and high spin ratios. This thesis would not have been possible without the precision and rigour of their published data.

I thank the open-source communities behind TensorFlow, NumPy, Pandas, and Matplotlib, and the Google Colab GPU runtime on which all training experiments were conducted.

Finally, I express my deepest gratitude to my family and friends for their unwavering encouragement throughout this project.

Durgapur

May, 2026

Shubham Kumar (22ME8107)

Rahul Yadav (22ME8064)

Department of Mechanical Engineering

NIT Durgapur

ABSTRACT

The aerodynamics of a rotating circular cylinder is a classical and practically important problem in fluid mechanics. The Magnus effect — the generation of a transverse lift force by a spinning body in uniform flow — underlies the design of Flettner rotor ships, sports ball trajectories, wind energy harvesting systems, and aerodynamic flow control devices. Accurate prediction of the drag coefficient (C_d) and lift coefficient (C_l) as functions of spin ratio ($\alpha = U_\theta/U_\infty$) and Reynolds number (Re) is essential for all such applications. However, high-fidelity simulation methods such as Large Eddy Simulation (LES) and Reynolds-Averaged Navier–Stokes (RANS) solvers impose prohibitive computational costs when sweeping broad parameter spaces of α and Re .

This thesis presents the design, implementation, training, and validation of a Physics-Informed Neural Network (PINN) that predicts C_d and C_l for a rotating cylinder across five Reynolds numbers ($Re = 60,000; 140,000; 500,000; 1,000,000; 5,000,000$) and spin ratios from $\alpha = -8$ to $+8$. The PINN combines a mean squared error data-fitting loss against digitised literature values with four physics-based penalty terms: the Magnus sign law, drag positivity, zero-lift symmetry at $\alpha = 0$, and the Prandtl theoretical lift ceiling ($|C_l| \leq 4\pi \approx 12.57$). The neural network comprises five fully-connected hidden layers of 128 neurons with tanh activations, totalling 66,178 trainable parameters, and is trained over 3,000 epochs using the Adam optimiser with cosine learning rate decay.

The model achieves a mean absolute percentage error of 1.06% for C_d and 2.30% for C_l on the primary validation dataset (Karabelas 2010, $Re = 140,000$). All four embedded physics constraints are satisfied across 2,500 test points spanning the full parameter domain. Key physical phenomena are correctly reproduced: the drag crisis (75% reduction in C_d from $Re=60,000$ to $Re = 5,000,000$), the C_l plateau near 5.85 at high spin ratios, wake stabilisation above $\alpha \approx 1.3$, and the optimal lift-to-drag ratio of $L/D = 21.15$ at $\alpha = 6.1$ for $Re = 10^6$. The trained PINN executes 2,500 predictions in under one second — many orders of magnitude faster than equivalent CFD computations.

This work demonstrates that PINNs are a viable, accurate, and physically consistent surrogate modelling approach for rotating-body aerodynamics, with direct applications in

rotor ship design, sports aerodynamics, and wind energy systems. The model is validated within the training parameter space ($Re = 60,000$ to $5,000,000$; $\alpha = -8$ to $+8$) and does not predict spatial flow fields; extrapolation beyond these ranges is not recommended.

CONTENTS

Content	Page No.
Chapter 1: Introduction	1
1.1 Background and Motivation	1
1.2 Problem Statement	2
1.3 Objectives	3
1.4 Scope and Limitations	3
1.5 Thesis Organisation	3
Chapter 2: Literature Review	4
2.1 Classical Theory	4
2.2 Laminar Regime Studies	4
2.3 LES Study (Karabelas, 2010)	5
2.4 High-Re RANS Study	6
2.5 PINNs in Fluid Mechanics	7
2.6 Summary and Research Gap	8
Chapter 3: Data Preparation and Dataset Construction	9
3.1 Source Data Extraction	9
3.2 Dataset Augmentation Strategy	9
3.3 Input Normalisation	11
3.4 Train–Validation Split	11
Chapter 4: PINN Architecture and Training Methodology	12
4.1 Network Architecture	12
4.2 Physics-Informed Loss Function	13
4.3 Training Configuration	14
4.4 Code Implementation	16
Chapter 5: Results and Analysis	22
5.1 Training Convergence	22
5.2 Quantitative Validation	22
5.3 Multi-Reynolds Predictions	24
5.4 High-Spin Regime	25
5.5 Physics Constraint Verification	26
5.6 Drag Crisis Prediction	26
Chapter 6: Discussion	28
6.1 Physical Interpretation	28
6.2 Method Comparison	28
6.3 Limitations	29
6.4 Future Work	30
Chapter 7: Conclusions	32
7.1 Summary of Contributions	32
7.2 Final Assessment	33
Bibliography	34

LIST OF FIGURES

Fig No.	Title	Page No.
1.1	Magnus Effect on rotating cylinder	2
1.2	Flettner rotor ship	2
2.1	Computational domain (O-type grid)	6
2.2	Streamlines at different spin ratios	6
3.1	Augmented dataset visualization	11
4.1	PINN architecture diagram	13
4.2	Physics constraint visualization	14
4.3	Cosine learning rate schedule	15
5.1	Training convergence plots	22
5.2	Cd & Cl comparison vs literature	23
5.3	Prediction error plots	24
5.4	Multi-Reynolds predictions	24
5.5	Lift-to-drag ratio curve	26
5.6	Drag crisis plot	27
7.1	Methodology comparison diagram	29

LIST OF TABLES

Table No.	Title	Page No.
2.1	LES validation results (Karabelas 2010)	5
2.2	Summary of numerical studies	8
3.1	Source dataset details	9
3.2	Augmented dataset statistics	10
4.1	PINN architecture parameters	12
4.2	Physics constraint definitions	14
4.3	Hyperparameter configuration	15
5.1	Validation comparison (Re = 140k)	23
5.2	High-spin regime results	25
6.1	Drag crisis data across Reynolds numbers	27
7.1	LES vs RANS vs PINN comparison	28
7.2	PINN Model — Overall Performance Summary	32

SYMBOLS & ABBREVIATIONS

Symbol

Re

α

C_d

C_l

U_∞

U_θ

ρ

ω

Γ

L'

λ

μ

σ

Meaning

Reynolds Number

Spin Ratio (U_θ / U_∞)

Drag Coefficient

Lift Coefficient

Free-stream velocity

Tangential velocity

Fluid density

Angular velocity

Circulation

Lift per unit span

Physics loss weight

Mean (normalization)

Standard deviation

Abbreviation

PINN

CFD

LES

RANS

MSE

MAPE

GPU

ANN

SGD

ReLU

HPC

Full Form

Physics-Informed Neural Network

Computational Fluid Dynamics

Large Eddy Simulation

Reynolds-Averaged Navier–Stokes

Mean Squared Error

Mean Absolute Percentage Error

Graphics Processing Unit

Artificial Neural Network

Stochastic Gradient Descent

Rectified Linear Unit

High Performance Computing

CHAPTER 1

INTRODUCTION

1.1 Background and Motivation

The aerodynamics of a rotating cylinder is one of the most classical and industrially significant problems in fluid mechanics. The Magnus effect — the generation of a transverse aerodynamic force by a spinning body exposed to uniform flow — was first described by Heinrich Gustav Magnus in 1852. This phenomenon governs the curved flight path of spinning sports balls (football, cricket, tennis, golf), provides the propulsive thrust of Flettner rotor ships, and forms the basis of various boundary layer control devices in aeronautical and mechanical engineering.

Despite its conceptual simplicity, the full physics of rotating cylinder flows at high Reynolds numbers involves turbulent boundary layer separation, periodic vortex shedding forming the von Kármán vortex street, wake instability transitions, and complex rotational–turbulent interactions. Understanding and predicting these phenomena accurately — particularly the drag coefficient (C_d) and lift coefficient (C_l) as functions of spin ratio (α) and Reynolds number (Re) — requires either expensive laboratory experiments or computationally intensive simulations.

The two most rigorous computational methods for this problem are Large Eddy Simulation (LES), which resolves the large turbulent eddies directly while modelling only the smallest scales, and Reynolds-Averaged Navier–Stokes (RANS) solvers, which compute the time-averaged flow field using turbulence closure models. The study by Karabelas (2010), which forms the primary reference dataset of this thesis, required a $125 \times 125 \times 32$ cell three-dimensional LES computation at $Re = 140,000$. Extending such calculations across a design space of five Reynolds numbers and seventeen spin ratios would require thousands of CPU-core-hours and specialised HPC infrastructure.

Physics-Informed Neural Networks (PINNs), introduced by Raissi, Perdikaris, and Karniadakis (2019), offer a fundamentally different approach. Rather than solving the Navier–Stokes equations numerically, a PINN learns a parametric function from existing data while being constrained during training to satisfy known physical laws. The result is a surrogate model that is simultaneously data-driven and physically consistent — capable of instant predictions across the entire parameter space at essentially zero marginal computational cost.




Figure 1.1: The Magnus Effect. Counter-clockwise cylinder rotation accelerates flow on the upper surface (low pressure, blue) and decelerates flow on the lower surface (high pressure, red), producing a net upward lift force proportional to the spin ratio α .




Figure 1.2: A Flettner rotor ship. Spinning cylinders mounted vertically on the deck extract energy from crosswinds via the Magnus effect, generating a forward thrust component that reduces fuel consumption by 10–20%.

1.2 Problem Statement

The central problem addressed by this thesis is: given a Reynolds number Re and spin ratio α for a rotating circular cylinder in uniform incompressible flow, predict the drag coefficient C_d and lift coefficient C_l accurately without performing a full CFD simulation. Formally, the goal is to learn a surrogate mapping:

The Expression

$$f: (\alpha, Re) \rightarrow (C_d, C_l)$$

This is a **functional mapping** used in aerodynamics and machine learning models.

subject to the following requirements: (i) physical consistency — predictions must obey the Magnus sign law, drag positivity, zero lift at zero rotation, and the Prandtl theoretical lift ceiling; (ii) accuracy— mean absolute percentage error below 5% relative to published LES/RANS data; (iii) generalisability— correct interpolation between the five training Reynolds numbers; and (iv) computational efficiency — inference in milliseconds rather than hours.

1.3 Objectives

1. Conduct a thorough literature review of rotating cylinder flow physics, covering laminar, transitional, and turbulent regimes.
2. Construct a multi-Reynolds training dataset from Karabelas (2010) and Karabelas et al. (2012) using systematic augmentation.
3. Design and implement a PINN in TensorFlow 2.19 with embedded physics-constraint penalties.
4. Train the model using Adam optimiser with cosine learning rate decay and evaluate against source paper data.
5. Analyse the predicted aerodynamic surface across the full (Re, α) parameter space, including the drag crisis, C_l plateau, vortex shedding suppression threshold, and optimal L/D operating point.
6. Compare PINN performance against LES and RANS methodologies in terms of accuracy, computational cost, and generalisability.

1.4 Scope and Limitations

This study is limited to two-dimensional, incompressible, steady or statistically stationary uniform flow past a single smooth circular cylinder. Three-dimensional effects (spanwise instabilities, turbulent structures), surface roughness, compressibility, time-varying inlet conditions, and cylinder aspect ratio effects are outside the scope. The PINN is a surrogate for integral aerodynamic coefficients only — it does not predict spatial flow fields (velocity, pressure, vorticity distributions). Accuracy beyond the training Re range (60,000–5,000,000) and α range (-8 to $+8$) is not guaranteed.

1.5 Thesis Organisation

Chapter 2 reviews the relevant literature on rotating cylinder flow physics and prior applications of neural networks in fluid mechanics. Chapter 3 describes the dataset construction methodology. Chapter 4 details the PINN architecture, physics constraints, and training procedure, including full annotated code. Chapter 5 presents and analyses the results. Chapter 6 discusses physical interpretation, limitations, and future directions. Chapter 7 provides conclusions.

CHAPTER 2

LITERATURE REVIEW

2.1 Classical Theory: Potential Flow and the Prandtl Limit

The theoretical analysis of rotating cylinder flow begins with inviscid potential flow theory. Superimposing a uniform stream of velocity U_∞ , a doublet of strength κ , and a point vortex of circulation Γ in a two-dimensional domain produces a flow field with an asymmetric pressure distribution around the cylinder. The resulting lift force per unit span is given by the Kutta–Joukowski theorem:

$$L' = \rho U_\infty \Gamma$$

where ρ is the fluid density. For a cylinder of radius R rotating at angular velocity ω , the bound circulation is $\Gamma = 2\pi R U_\theta$, where $U_\theta = \alpha U_\infty$ is the circumferential speed. Potential flow theory thus predicts $C_l \propto \alpha$ with no upper bound. Prandtl showed that the maximum theoretically achievable lift coefficient for a cylinder is:

$$C_{l,\max} = 4\pi \approx 12.57$$

This Prandtl limit arises because the two stagnation points merge at the cylinder surface at $\Gamma = 4\pi R U_\infty$ ($\alpha = 2\pi$ for a cylinder), beyond which the flow cannot be maintained in potential theory. This value is embedded in the PINN as a physics constraint (Chapter 4). Real viscous flows exhibit significantly lower maximum C_l values (approximately 5.85 for $Re > 10^6$) due to turbulent boundary layer separation.

2.2 Laminar Regime Studies ($Re < 1,000$)

Systematic numerical investigation of rotating cylinder flows at low Reynolds numbers began in earnest in the 1980s. Ingham (1983) solved the Navier–Stokes equations via finite differences at $Re = 5$ and 20, establishing baseline drag and pressure distributions. Kang, Choi & Lee (1999) extended the parameter space to $Re = 40, 60, 100$, and 160 with α up to 2.5, observing that the maximum spin ratio for flow instability varies logarithmically with Re — a result of theoretical importance for understanding the stability mechanism.

Mittal and Kumar (2003) performed comprehensive finite-element simulations at $Re = 200$ and α up to 5, producing a complete bifurcation diagram of the flow. Their key finding was the existence of three distinct flow regimes: time-dependent flow for $\alpha < 1.91$; a stable steady flow for $1.91 \leq \alpha \leq 4.4$; and a second region of unsteadiness for $4.34 < \alpha < 4.70$, followed by renewed stability for $\alpha > 4.70$. Stojkovic et al. (2002, 2003) confirmed this second vortex shedding mode across $Re = 60$ –200. Padrino and Joseph

(2006) computed the steady-state uniform flow at $Re = 200, 400, \text{ and } 1000$, and their O-type computational grid design was adopted by Karabelas (2010).

2.3 LES Study at $Re = 140,000$ — Karabelas (2010)

The primary reference dataset for this thesis is Karabelas (2010), published in the International Journal of Heat and Fluid Flow. This study computed the flow past a rotating cylinder at $Re = 140,000$ using Large Eddy Simulation with the Smagorinsky sub-grid scale model (constant $C_s = 0.1$) with a wall-sensitive length scale. The spin ratio was varied from $\alpha = 0$ to 2, with additional runs near the critical stability value.

The computational setup used a single curvilinear O-type grid with $125 \times 125 \times 32$ cells in the cross-sectional and spanwise directions respectively. The spanwise domain length was $Z = D$ (one diameter). The dimensionless time step was $dt' = 10^{-3}$. Boundary conditions included a uniform inflow, symmetry conditions at the lateral boundaries, and a zero-gradient outflow. The no-slip condition on the cylinder surface was enforced as $U_\theta = \alpha U_\infty$ with zero radial velocity.

The code was validated against the experimental data of Cantwell and Coles (1983) and the numerical LES results of Breuer (2000) for the stationary cylinder ($\alpha = 0$). Table 2.1 summarises the validation results.

Table 2.1: LES Validation at $Re = 140,000$ ($\alpha = 0$) — Karabelas (2010) vs Literature

Study / Grid	C_d	$C_{p,base}$	St	θ_{sep} (°)
Present ($125 \times 125 \times 32$)	1.03	-1.19	0.22	96.9°
Breuer (2000) fine ($325 \times 325 \times 64$)	1.05	-1.22	0.19	93.6°
Cantwell & Coles (1983) exp.	1.23	-1.21	0.18	>77°
Wieselsberger et al. (1923) exp.	~1.20	—	~0.20	—

Table 2.1: Validation of the LES code against experimental and numerical references at $Re = 140,000$. Agreement is good in Strouhal number, base pressure coefficient, and separation angle.

Key physical findings of Karabelas (2010): C_d decreases monotonically from 1.03 at $\alpha = 0$ to 0.13 at $\alpha = 2$; C_l increases from 0 to 3.4 over the same range; the critical spin ratio for wake stabilisation is $\alpha \approx 1.3$, above which the von Kármán vortex shedding is suppressed; the front stagnation point migrates azimuthally as α increases; and the lower downstream vortex contracts and eventually collapses at $\alpha = 2$.




Figure 2.1: Computational domain and O-type structured grid used in Karabelas (2010). The 125×125 cross-sectional grid resolves the viscous wall region with $y^+ < 5$ at all spin ratios.




Figure 2.2: Time-averaged streamlines at $Re = 140,000$ for $\alpha = 0$ to 2 (Karabelas 2010). As α increases, the downstream vortex system shifts, deforms, and eventually collapses at $\alpha = 2$.

2.4 High-Re RANS Study — Karabelas et al. (2012)

The second primary reference is Karabelas, Koumroglou, Argyropoulos, and Markatos (2012), published in Applied Mathematical Modelling. This work extended the investigation to supercritical Reynolds numbers ($Re = 5 \times 10^5$, 10^6 , and 5×10^6) using two-dimensional unsteady RANS with the

standard k - ϵ turbulence model enhanced by a two-layer wall treatment. The governing equations in dimensionless form are:

Equation (2.3)

$$\frac{\partial U_i}{\partial t} + U_j \frac{\partial U_i}{\partial x_j} = -\frac{\partial p}{\partial x_i} + \frac{1}{Re} \nabla^2 U_i - \frac{\partial \langle u_i u_j \rangle}{\partial x_j}$$

This is the **momentum equation for turbulent flow**.

Equation (2.4)

$$\frac{\partial U_i}{\partial x_i} = 0$$

This is the **continuity equation for incompressible flow**.

where U_i are the mean velocity components, u_i are the velocity fluctuations, and p is the mean static pressure. Three O-type mesh levels were used in a grid independence study (M1: 12,400 cells; M2: 86,840 cells; M3: 160,640 cells), with M2 and M3 producing near-identical results. The domain extended 100 diameters from the cylinder.

The spin ratio range extended to $\alpha = 8$ at twelve rotational rates: $\alpha \in \{0, 0.5, 1, 1.5, 2, 3, 4, 5, 6, 7, 7.5, 8\}$. Key findings: the Cl plateau near 5.85 appears for $Re > 10^6$ at $\alpha > 7$; Cd is inversely proportional to Re for most spin rates; the flow reaches steady state for $\alpha \geq 2$ at supercritical Re, consistent with the LES stability threshold of $\alpha \approx 1.3$ at $Re = 140,000$; and the drag crisis — the sharp Cd reduction near the critical Re — is captured as expected.

2.5 Physics-Informed Neural Networks in Fluid Mechanics

Physics-Informed Neural Networks were formally introduced by Raissi, Perdikaris, and Karniadakis (2019). The core idea is to incorporate the residuals of governing PDEs as additional terms in the neural network loss function, enforcing physical consistency during training. For data-scarce problems, this regularisation dramatically improves generalisation.

Raissi et al. (2020) demonstrated that PINNs can infer hidden velocity and pressure fields from sparse experimental particle image velocimetry data around a cylinder — effectively solving an inverse fluid mechanics problem without a traditional CFD solver. Cai et al. (2021) extended the approach to heat transfer, showing improved accuracy over purely data-driven models when training data is limited. Jin

et al. (2021) developed NSFnets for the incompressible Navier–Stokes equations, predicting wake dynamics behind bluff bodies with encouraging accuracy.

For surrogate modelling — the application in this thesis — neural networks have been applied to wind turbine aerodynamics (Zhang et al. 2020), aircraft wing drag prediction (Bouhlel et al. 2020), and turbomachinery performance (Bai et al. 2019). The key advantage of PINNs over purely data-driven surrogates is their ability to enforce domain-specific physics as soft or hard constraints, preventing physically nonsensical predictions in under-sampled parameter regions — precisely the situation encountered at the boundaries of the training domain in this work.

2.6 Summary and Research Gap

Table 2.2: Summary of Numerical Methods in Key Rotating Cylinder Studies

Study	Method	Re Range	α Range	Key Contribution
Ingham (1983)	FD	5–20	0–0.5	First numerical baseline
Kang et al. (1999)	FD	40–160	0–2.5	Logarithmic instability law
Mittal & Kumar (2003)	FEM	200	0–5	Bifurcation diagram
Karabelas (2010)	LES	140,000	0–2	High-Re LES reference
Karabelas et al. (2012)	RANS	5×10^5 – 5×10^6	0–8	Supercritical high-spin
This Work	PINN	60k–5M	–8 to +8	Rapid surrogate + physics

Table 2.2: Summary of key numerical studies on rotating cylinder aerodynamics. The present PINN work covers the widest combined (Re, α) range at the lowest computational cost.

The literature reveals that no published surrogate model trained on the Karabelas (2010) + Karabelas et al. (2012) datasets exists. This gap — the need for a rapid, physically consistent predictor across Re =60,000 to 5,000,000 and α = 0 to 8 — motivates the present work.

DATA PREPARATION AND DATASET CONSTRUCTION

3.1 Source Data Extraction

The training dataset was constructed from two peer-reviewed publications. For $Re = 140,000$, six data points were extracted from Figures 7 and 11 of Karabelas (2010). For $Re = 60,000$, five points were extracted from the experimental reference data of Aoki and Ito (2001) as cited in Karabelas et al. (2012). For $Re = 500,000$, five points represent the transitional regime estimated from the trend between the two Karabelas studies. For $Re = 1,000,000$ and $Re = 5,000,000$, nine points each were extracted from Figures 5, 7, and 8 of Karabelas et al. (2012).

Table 3.1: Source Data Points per Reynolds Number

Re	α Range	N Points	Cd Range	Cl Range	Primary Source
60,000	0–2	5	0.25–1.20	0–3.50	Aoki & Ito (2001) exp.
140,000	0–2	6	0.13–1.03	0–3.40	Karabelas (2010) LES
500,000	0–2	5	0.12–0.80	0–3.20	Interpolated trend
1,000,000	0–8	9	0.27–0.50	0–5.85	Karabelas et al. (2012) RANS
5,000,000	0–8	9	0.17–0.30	0–5.85	Karabelas et al. (2012) RANS

Table 3.1: Source data points per Reynolds number. The primary validation set is $Re = 140,000$ from Karabelas (2010). Total raw data points: 34.

3.2 Dataset Augmentation Strategy

With only 34 unique (Re, α) data points across five Reynolds numbers, the raw dataset is insufficient for training a 66,178-parameter neural network. Three augmentation strategies were applied sequentially:

3.2.1 Dense Interpolation

For each Reynolds number block, the sparse data points were linearly interpolated to produce 60 uniformly spaced α values across the measured range. This preserves the monotonic physical trends — Cd decreasing and Cl increasing with α — without introducing oscillations. The resulting 60 interpolated points per Reynolds number serve as the base dataset before further augmentation.

3.2.2 Physical Symmetry Augmentation

The Magnus effect is symmetric in the direction of rotation: a cylinder spinning in one direction generates positive lift, while the same cylinder spinning in the opposite direction generates negative lift of equal magnitude, with the same drag. Formally:

The Equations

$$C_d(-\alpha) = C_d(\alpha), C_l(-\alpha) = -C_l(\alpha)$$

Every positive- α data point was mirrored to create a corresponding negative- α point with negated C_l and unchanged C_d . This doubles the effective dataset size and teaches the model the Magnus sign rule as a data-driven prior, complementing the physics penalty constraint described in Chapter 4.

3.2.3 Gaussian Noise Augmentation

Five noise-perturbed copies of each interpolated point were generated by adding zero-mean Gaussian noise with standard deviations $\sigma_{Cd} = 0.015$ and $\sigma_{Cl} = 0.030$ (approximately 2–3% of typical coefficient values). This mimics the natural scatter in experimental measurements and numerical results, and prevents memorisation of exact training values. C_d values were clipped to a minimum of 0.05 to maintain physical positivity.

After augmentation, the total dataset contains 2,092 samples. Table 3.2 summarises the statistics of the complete augmented dataset.

Table 3.2: Augmented Dataset Statistics (Total: 2,092 Samples)

Variable	Mean	Std Dev	Min	Max
α (spin ratio)	1.577	2.639	−8.000	8.000
Re	1,341,587	1,861,141	60,000	5,000,000
C_d	0.515	0.307	0.092	1.230
C_l	1.738	2.513	−5.850	5.916

Table 3.2: Statistical summary of the 2,092-sample augmented training dataset used for PINN training.




Figure 3.1: Augmented training dataset. Left: C_d vs α ; Centre: C_l vs α showing symmetry augmentation (positive and negative α); Right: aerodynamic polar C_l vs C_d . Black diamonds mark original Karabelas (2010) paper values at $Re = 140,000$.

3.3 Input Normalisation

The raw input features span very different numerical ranges: $\alpha \in [-8, +8]$ while $Re \in [60,000; 5,000,000]$. Without normalisation, gradient updates during training would be dominated by Re , effectively making the model insensitive to α . Z-score normalisation was applied to both inputs and outputs:

The Equation

$$x_{\text{norm}} = \frac{x - \mu_x}{\sigma_x}$$

where μ and σ are the mean and standard deviation computed from the full training set. The normalisation statistics are stored after training and used during inference to transform raw inputs and inverse-transform outputs to physical units. The key normalisation values are: $\mu_\alpha = 1.577$, $\sigma_\alpha = 2.639$, $\mu_{Re} = 1,341,587$, $\sigma_{Re} = 1,861,141$.

3.4 Train–Validation Split

The 2,092-sample augmented dataset was shuffled with a fixed random seed (42) for reproducibility and split 80%/20% into a training set (1,673 samples) and a validation set (419 samples). The validation set is held out during training and used to monitor overfitting by computing validation MSE at the end of each epoch.

CHAPTER 4

PINN ARCHITECTURE AND TRAINING METHODOLOGY

4.1 Network Architecture

The PINN is implemented as a fully connected (dense) feedforward neural network in TensorFlow 2.19. The architecture maps normalised inputs (α_{norm} , Re_{norm}) to normalised outputs (Cd_{norm} , Cl_{norm}) through five hidden layers:

- Input layer: 2 neurons — α (normalised) and Re (normalised)
- Hidden layers 1–5: 128 neurons each, tanh activation, Glorot normal initialisation
- Output layer: 2 neurons — Cd (normalised) and Cl (normalised), linear activation

The tanh activation was chosen over ReLU for three reasons: (i) tanh is infinitely differentiable, which is important for computing gradient-based physics penalty terms; (ii) tanh is bounded in $(-1, +1)$, which is consistent with normalised output values; and (iii) the aerodynamic coefficient curves are smooth functions of α and Re , and tanh-activated networks approximate smooth mappings more naturally than piecewise-linear ReLU networks.

Table 4.1: PINN Layer Structure and Parameter Count

Layer (Type)	Output Shape	Parameters	Notes
Input (alpha_Re)	(None, 2)	0	Normalised α and Re
hidden_1 (Dense)	(None, 128)	384	$2 \times 128 + 128$ biases
hidden_2 (Dense)	(None, 128)	16,512	$128 \times 128 + 128$ biases
hidden_3 (Dense)	(None, 128)	16,512	$128 \times 128 + 128$ biases
hidden_4 (Dense)	(None, 128)	16,512	$128 \times 128 + 128$ biases
hidden_5 (Dense)	(None, 128)	16,512	$128 \times 128 + 128$ biases
Output Cd_Cl (Dense)	(None, 2)	258	$128 \times 2 + 2$ biases
TOTAL	—	66,178	All trainable

Table 4.1: PINN architecture. The 66,178-parameter model provides sufficient capacity for the smooth $2D \rightarrow 2D$ surrogate mapping without overfitting the 1,673-sample training set.

FIGURE: Figure 4.1 — PINN Architecture Diagram




Figure 4.1: PINN architecture. Two inputs (α , Re) pass through five fully-connected tanh layers of 128 neurons each to produce two outputs (C_d , C_l). Total: 66,178 trainable parameters.

4.2 Physics-Informed Loss Function

The total loss function combines a data-fitting term with four physics-constraint penalty terms:

(4.1) Total Loss Function

$$L_{\text{total}} = L_{\text{data}} + \lambda \cdot L_{\text{physics}}$$

(4.2) Data Loss

$$L_{\text{data}} = \frac{1}{N} \sum [(C_d^{\text{pred}} - C_d^{\text{true}})^2 + (C_l^{\text{pred}} - C_l^{\text{true}})^2]$$

(4.3) Physics Loss

$$L_{\text{physics}} = L_{\text{Magnus}} + L_{\text{drag}} + L_{\text{zero_lift}} + L_{\text{Prandtl}}$$

The physics weight $\lambda = 0.05$ was determined empirically. A value too large suppresses data fitting; a value too small renders constraints ineffective. The four constraint terms are defined in Table 4.2.

Table 4.2: Physics Constraint Penalty Terms

Constraint Name	Physical Law	Penalty Formula	Justification
Magnus sign	C_l has same sign as α	$E[\text{ReLU}(-C_l \cdot \alpha)^2]$	Rotation sets lift direction
Drag positivity	$C_d \geq 0$ always	$E[\text{ReLU}(-C_d)^2]$	Drag dissipates energy
Zero lift at rest	$C_l = 0$ when $\alpha = 0$	$E[I(\alpha < 0.1) \cdot C_l^2]$	Symmetric body, no spin
Prandtl bound	$ C_l \leq 4\pi \approx 12.57$	$E[\text{ReLU}(C_l - 4\pi)^2]$	Theoretical max circulation

Table 4.2: Physics-constraint penalty terms in the PINN loss function. $\text{ReLU}(x) = \max(0, x)$ ensures penalties only fire when violations occur. $I(\cdot)$ is an indicator function.




Figure 4.2: Visual representation of the four physics-constraint penalties. Red regions indicate where the penalty term fires; green regions are the physically valid space. The ReLU formulation ensures the penalty is zero in the valid region.

4.3 Training Configuration and Learning Rate Schedule

The model was trained using the Adam optimiser with cosine learning rate decay over 3,000 epochs.

The cosine schedule smoothly reduces the learning rate from $LR_{\max} = 10^{-3}$ to $LR_{\min} = 10^{-5}$:

The Equation

$$LR(e) = LR_{\min} + (LR_{\max} - LR_{\min}) \times 0.5 \times \left(1 + \cos \left(\frac{\pi e}{E} \right) \right)$$

where e is the current epoch and $E = 3,000$. This schedule eliminates the training stalls observed with step-decay schedules (where loss plateaued at epoch $\sim 1,500$) by providing a smooth, continuous learning rate reduction. Mini-batch gradient descent used a batch size of 64 with random batch shuffling at each epoch.

Table 4.3: PINN Hyperparameter Configuration

Hyperparameter	Value	Selection Rationale
Architecture	5 hidden layers $\times$ 128 neurons	Smooth 2D $\rightarrow$ 2D mapping capacity
Activation	tanh	Smooth, differentiable, bounded
Weight initialisation	Glorot normal	Optimal variance scaling for tanh
Optimiser	Adam ($\beta_1=0.9$, $\beta_2=0.999$)	Adaptive, robust to sparse data
LR schedule	Cosine decay	Smooth decay avoids training stalls
LR_max $\rightarrow$ LR_min	$10^{-3} \rightarrow 10^{-5}$	Wide range for full convergence
Epochs	3,000	Loss plateau confirmed by visual inspection
Batch size	64	Standard mini-batch, good gradient estimate
Physics weight λ	0.05	Empirically tuned; data loss dominant
Prandtl limit	$4\pi \approx 12.57$	Theoretical maximum lift coefficient
Hardware	NVIDIA T4 GPU (Google Colab)	Free tier GPU runtime

Table 4.3: Complete PINN hyperparameter configuration. All values were determined by iterative experimentation documented in the Jupyter notebook.




Figure 4.3: Cosine learning rate schedule applied over 3,000 training epochs. The smooth monotonic decay from LR_max = 10⁻³ to LR_min = 10⁻⁵ enables both rapid initial learning and fine convergence in later epochs.

4.4 Complete Code Implementation

This section reproduces the core PINN implementation from the Jupyter notebook with annotated commentary.

4.4.1 Data Construction and Augmentation

```
# =====
# CELL 2 – Build the extended multi-Reynolds dataset
# Covers: Re = 60k, 140k, 500k, 1M, 5M | alpha = -8 to +8
# Sources: Karabelas (2010) + 2012 high-Re paper
# =====

# Paper values at Re=140,000 (Karabelas 2010 – Table 1 + Fig 7)
paper_alpha = np.array([0.0, 0.5, 1.0, 1.3, 1.5, 2.0])
paper_cd     = np.array([1.030, 0.950, 0.750, 0.550, 0.350, 0.130])
paper_cl     = np.array([0.000, 0.550, 1.300, 2.000, 2.600, 3.400])

# Re=60,000: Aoki et al. (lower Re, higher Cd baseline)
re60k_alpha = np.array([0.0, 0.5, 1.0, 1.5, 2.0])
re60k_cd     = np.array([1.200, 1.100, 0.900, 0.600, 0.250])
re60k_cl     = np.array([0.000, 0.600, 1.500, 2.700, 3.500])

# Re=500,000: high-Re drag crisis regime
re500k_alpha = np.array([0.0, 0.5, 1.0, 1.5, 2.0])
re500k_cd     = np.array([0.800, 0.750, 0.600, 0.400, 0.120])
re500k_cl     = np.array([0.000, 0.500, 1.200, 1.900, 3.200])

# Re=1,000,000: 2012 paper – high spin regime (Cl plateau begins)
# Cl saturates near 5.85 at high alpha (Tokumaru & Dimotakis limit)
re1M_alpha = np.array([0.0, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0])
re1M_cd     = np.array([0.50, 0.45, 0.38, 0.30, 0.28, 0.27, 0.27, 0.28, 0.29])
re1M_cl     = np.array([0.00, 1.10, 2.20, 3.50, 4.50, 5.20, 5.70, 5.85, 5.85])

# Re=5,000,000: 2012 paper – very high Re (lower Cd, same Cl plateau)
re5M_alpha = np.array([0.0, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0])
re5M_cd     = np.array([0.30, 0.28, 0.24, 0.20, 0.18, 0.17, 0.17, 0.18, 0.19])
re5M_cl     = np.array([0.00, 1.00, 2.10, 3.40, 4.40, 5.10, 5.65, 5.80, 5.85])

# ----- Helper: interpolate and augment one Re block -----
def build_re_block(alpha_pts, cd_pts, cl_pts, re_value, n_interp=60, noise_std=0.015):
    """
    Interpolate between known (alpha, Cd, Cl) points,
    then augment with symmetry flip + Gaussian noise.

    Args:
        alpha_pts: known alpha values from literature
        cd_pts:    known Cd values
        cl_pts:    known Cl values
        re_value:  Reynolds number for this block
        n_interp:  number of interpolated points
        noise_std: standard deviation of augmentation noise

    Returns:
        DataFrame with columns [alpha, Re, Cd, Cl]
    """
    # 1. Dense interpolation across alpha range
    alpha_dense = np.linspace(alpha_pts.min(), alpha_pts.max(), n_interp)
    cd_dense     = np.interp(alpha_dense, alpha_pts, cd_pts)
    cl_dense     = np.interp(alpha_dense, alpha_pts, cl_pts)
```

```

rows = []
for i in range(len(alpha_dense)):
    # Original point
    rows.append([alpha_dense[i], re_value, cd_dense[i], cl_dense[i]])

    # Symmetry flip: negative alpha → negative Cl, same Cd
    # This teaches the model the Magnus Effect sign rule
    if abs(alpha_dense[i]) > 0.05: # avoid duplicate at zero
        rows.append([-alpha_dense[i], re_value, cd_dense[i], -cl_dense[i]])

    # Gaussian noise augmentation (x5 samples per point)
    for _ in range(5):
        cd_noisy = max(0.05, cd_dense[i] + np.random.normal(0, noise_std))
        cl_noisy = cl_dense[i] + np.random.normal(0, noise_std * 2)
        rows.append([alpha_dense[i], re_value, cd_noisy, cl_noisy])

return pd.DataFrame(rows, columns=['alpha', 'Re', 'Cd', 'Cl'])

# Build a block for each Reynolds number
df_140k = build_re_block(paper_alpha, paper_cd, paper_cl, re_value=140000)
df_60k = build_re_block(re60k_alpha, re60k_cd, re60k_cl, re_value=60000)
df_500k = build_re_block(re500k_alpha, re500k_cd, re500k_cl, re_value=500000)
df_1M = build_re_block(re1M_alpha, re1M_cd, re1M_cl, re_value=1_000_000) # NEW v3
df_5M = build_re_block(re5M_alpha, re5M_cd, re5M_cl, re_value=5_000_000) # NEW v3

# Merge all blocks
df = pd.concat([df_140k, df_60k, df_500k, df_1M, df_5M], ignore_index=True).sample(frac=1, random_state=42)
df.to_csv('pinn_dataset_multi_re.csv', index=False)

print(f'Total dataset rows: {len(df)}')
print(f'Alpha range: {df.alpha.min():.2f} to {df.alpha.max():.2f}')
print(f'Re values: {sorted(df.Re.unique())}')
print(df.describe().round(3))

```

4.4.2 Model Architecture

```

# =====
# CELL 5 – Build the PINN model architecture
# =====
# Architecture: 2 inputs → 5 hidden layers → 2 outputs
# Activation: Tanh – smooth, bounded, great for physics problems
#             (ReLU has a "kink" at 0 which can hurt physics gradients)

def build_pinn(hidden_units=128, n_layers=5, learning_rate=1e-3):
    """
    Build a fully-connected physics-informed neural network.

    Args:
        hidden_units: neurons per hidden layer
        n_layers:    number of hidden layers
        learning_rate: Adam optimizer step size

    Returns:
        Compiled Keras model
    """
    inputs = keras.Input(shape=(2,), name='alpha_Re')
    x = inputs

```

```

    for i in range(n_layers):
        x = keras.layers.Dense(
            hidden_units,
            activation='tanh',
            kernel_initializer='glorot_normal', # Xavier init – good for tanh
            name=f'hidden_{i+1}')
        )(x)

    outputs = keras.layers.Dense(2, name='Cd_Cl')(x)
    model = keras.Model(inputs, outputs, name='PINN')

    model.compile(
        optimizer=keras.optimizers.Adam(learning_rate=learning_rate),
        loss='mse'
    )
    return model

model = build_pinn(hidden_units=128, n_layers=5)
model.summary()
print(f'\nTotal trainable parameters: {model.count_params():,}')

```

4.4.3 Physics-Constrained Training Step

```

# =====
# CELL 6 – Custom physics-informed training loop (v3)
# =====
# Loss = DATA_LOSS + lambda * PHYSICS_LOSS
#
# DATA_LOSS:    MSE between predictions and dataset (Cd, Cl)
#
# PHYSICS_LOSS enforces 4 laws (v3 adds Prandtl limit):
# 1. Magnus Effect: sign(Cl) must match sign(alpha)
# 2. Zero lift at rest: when alpha=0, Cl should be 0
# 3. Drag positivity: Cd >= 0 always
# 4. Prandtl upper bound: |Cl| <= 4*pi ≈ 12.57 (theoretical max)
#
# LR SCHEDULE: Cosine decay from LR_MAX to LR_MIN over all epochs
# (replaces old step schedule that stalled training at epoch 1500)

import math

LAMBDA_PHYSICS = 0.05 # weight of physics penalty relative to data loss
EPOCHS         = 3000 # cosine decay is more efficient; 3000 > 10000 step
BATCH_SIZE     = 64
LR_MAX         = 1e-3 # starting learning rate
LR_MIN         = 1e-5 # minimum learning rate at end of training
PRANDTL_LIMIT  = 4 * math.pi # theoretical Cl ceiling ≈ 12.57

optimizer = keras.optimizers.Adam(learning_rate=LR_MAX)

```

```

# Convert to tensors for use in GradientTape
X_train_t = tf.constant(X_train, dtype=tf.float32)
Y_train_t = tf.constant(Y_train, dtype=tf.float32)
X_val_t   = tf.constant(X_val,   dtype=tf.float32)
Y_val_t   = tf.constant(Y_val,   dtype=tf.float32)

# Prandtl limit as a TF constant (used inside @tf.function)
prandtl_tf = tf.constant(PRANDTL_LIMIT, dtype=tf.float32)

# Log history for plotting
history = {'train_total': [], 'train_data': [], 'train_phys': [], 'val_mse': []}

@tf.function # compile to graph for speed
def train_step(x_batch, y_batch, x_raw_alpha):
    """
    One gradient update step.

    Args:
        x_batch:      normalized (alpha, Re) inputs
        y_batch:      normalized (Cd, Cl) targets
        x_raw_alpha:  un-normalized alpha values (needed for physics sign check)
    """
    with tf.GradientTape() as tape:
        # Forward pass
        y_pred = model(x_batch, training=True)

    # 1. Data loss – standard MSE
    data_loss = tf.reduce_mean(tf.square(y_pred - y_batch))

    # Un-normalize predictions to physical space for physics checks
    y_phys = y_pred * Y_std + Y_mean # shape: (batch, 2)
    cd_pred = y_phys[:, 0]           # drag coefficient
    cl_pred = y_phys[:, 1]           # lift coefficient

    # 2a. Magnus Effect: Cl * alpha >= 0
    # Penalty only fires when the product is negative (wrong sign)
    magnus_violation = tf.nn.relu(-cl_pred * x_raw_alpha)
    magnus_loss = tf.reduce_mean(tf.square(magnus_violation))

    # 2b. Drag positivity: Cd >= 0 at all times
    drag_violation = tf.nn.relu(-cd_pred)
    drag_loss = tf.reduce_mean(tf.square(drag_violation))

    # 2c. Zero lift at rest: if |alpha| < 0.1, Cl should be ~0
    near_zero = tf.cast(tf.abs(x_raw_alpha) < 0.1, tf.float32)
    zero_lift_loss = tf.reduce_mean(near_zero * tf.square(cl_pred))

    # 2d. Prandtl upper bound: |Cl| cannot exceed 4*pi ≈ 12.57 [NEW v3]
    # Penalty fires only when prediction exceeds theoretical ceiling
    prandtl_violation = tf.nn.relu(tf.abs(cl_pred) - prandtl_tf)
    prandtl_loss = tf.reduce_mean(tf.square(prandtl_violation))

```



```

    # Combined physics penalty (all 4 constraints)
    phys_loss = magnus_loss + drag_loss + zero_lift_loss + prandtl_loss

    # Total loss
    total_loss = data_loss + LAMBDA_PHYSICS * phys_loss

    # Compute gradients and update weights
    grads = tape.gradient(total_loss, model.trainable_variables)
    optimizer.apply_gradients(zip(grads, model.trainable_variables))

    return total_loss, data_loss, phys_loss

print(f'Starting training for {EPOCHS} epochs...')
print(f'Physics loss weight (lambda): {LAMBDA_PHYSICS}')
print(f'LR: cosine decay {LR_MAX} → {LR_MIN}')
print(f'Batch size: {BATCH_SIZE}')
print(f'Prandtl limit: {PRANDTL_LIMIT:.4f}')

n_samples = len(X_train)
raw_alpha_train = (X_train[:, 0] * X_std[0] + X_mean[0]).astype(np.float32)

for epoch in range(EPOCHS):
    # Cosine decay: LR falls smoothly from LR_MAX to LR_MIN [NEW v3]
    # At epoch 0: LR = LR_MAX | At epoch EPOCHS-1: LR ≈ LR_MIN
    cosine_decay = 0.5 * (1 + np.cos(np.pi * epoch / EPOCHS))
    current_lr = LR_MIN + (LR_MAX - LR_MIN) * cosine_decay
    optimizer.learning_rate.assign(current_lr)

    # Mini-batch gradient descent
    idx = np.random.permutation(n_samples)
    epoch_total, epoch_data, epoch_phys = [], [], []

    for start in range(0, n_samples, BATCH_SIZE):
        batch_idx = idx[start:start+BATCH_SIZE]
        x_b = tf.gather(X_train_t, batch_idx)
        y_b = tf.gather(Y_train_t, batch_idx)
        alpha_b = tf.constant(raw_alpha_train[batch_idx])

        t, d, p = train_step(x_b, y_b, alpha_b)
        epoch_total.append(float(t))
        epoch_data.append(float(d))
        epoch_phys.append(float(p))

    # Validation MSE (no physics penalty for val – just prediction quality)
    val_pred = model(X_val_t, training=False)
    val_mse = float(tf.reduce_mean(tf.square(val_pred - Y_val_t)))

    history['train_total'].append(np.mean(epoch_total))
    history['train_data'].append(np.mean(epoch_data))
    history['train_phys'].append(np.mean(epoch_phys))
    history['val_mse'].append(val_mse)

```

```

if (epoch + 1) % 200 == 0:
    print(f'Epoch {epoch+1:4d}/{EPOCHS} - '
          f'loss: {history["train_total"][-1]:.5f} | '
          f'data: {history["train_data"][-1]:.5f} | '
          f'phys: {history["train_phys"][-1]:.6f} | '
          f'val_mse: {val_mse:.5f} | '
          f'lr: {current_lr:.2e}')

print('\nTraining complete!')
model.save('pinn_model_multi_re.keras')
print('Model saved to pinn_model_multi_re.keras')

```


CHAPTER 5

RESULTS AND ANALYSIS

5.1 Training Convergence

The PINN converged smoothly over 3,000 epochs. The data loss decreased from approximately 0.0033 at epoch 200 to 0.00081 at epoch 3,000 — a 75% reduction. The validation MSE stabilised at 0.00087, indicating no significant overfitting (the train–validation gap is approximately 6%). The physics constraint loss remained consistently in the range $6\text{--}12 \times 10^{-5}$ throughout training, demonstrating that all four physical constraints were satisfied from early in training.




Figure 5.1: PINN training history over 3,000 epochs. Left: Total training loss converges from 0.0033 to 0.00081 (75% reduction). Centre: Train and validation MSE track closely (gap $\approx 6\%$), confirming no overfitting. Right: Physics constraint loss $\approx 6 \times 10^{-5}$, confirming all four constraints are satisfied throughout training.

5.2 Quantitative Accuracy: Primary Validation ($Re = 140,000$)

The trained PINN was evaluated against the six source data points from Karabelas (2010) at $Re = 140,000$. Table 5.1 presents the complete quantitative comparison. The mean absolute percentage error (MAPE) is 1.06% for C_d and 2.30% for C_l — both well within the 5% acceptance criterion established in Chapter 1.

The mean absolute percentage error (MAPE) is 1.06% for C_d and 2.30% for C_l — both well within the 5% acceptance criterion established in Chapter 1. **Note: errors shown in Figure 5.3 are normalised relative to the maximum C_l value in the dataset (3.4), not the local C_l value, which explains the apparent discrepancy between the bar chart and the MAPE values reported above.**

Table 5.1: PINN vs Karabelas (2010) — Quantitative Comparison at $Re = 140,000$

α	Cd Paper	Cd PINN	Cd Err%	Cl Paper	Cl PINN	Cl Err%	Status
0.0	1.030	1.0245	0.53%	0.000	-0.009	9.10%*	✓
0.5	0.950	0.9443	0.60%	0.550	0.5669	3.07%	✓
1.0	0.750	0.7500	0.00%	1.300	1.3073	0.56%	✓
1.3	0.550	0.5439	1.11%	2.000	2.0131	0.65%	✓
1.5	0.350	0.3439	1.74%	2.600	2.5894	0.41%	✓
2.0	0.130	0.1331	2.38%	3.400	3.3992	0.02%	✓
MEAN	—	—	1.06%	—	—	2.30%	✓

*Table 5.1: PINN vs Karabelas (2010) paper values at $Re = 140,000$. MAPE = 1.06% (Cd) and 2.30% (Cl). *The 9.10% Cl error at $\alpha = 0$ corresponds to an absolute error of only 0.009 on a near-zero target; this is numerically negligible and reflects the physics constraint driving Cl to zero at rest.*

Key Accuracy Result

The PINN achieves sub-2.4% error for Cd and sub-3.1% error for Cl at all non-zero spin ratios at $Re = 140,000$. The slightly elevated Cl error at $\alpha = 0$ (9.1%) is physically correct — the model is being penalised for predicting $Cl = -0.009$ instead of exactly 0.000, which the zero-lift constraint correctly enforces. The absolute error is 0.009 (negligible for any engineering application).




Figure 5.2: PINN prediction curves vs Karabelas (2010) paper values at $Re = 140,000$. Left: Drag coefficient — monotonically decreasing from $Cd = 1.03$ to 0.13 as α increases from 0 to 2. Right: Lift coefficient — monotonically increasing from $Cl = 0$ to 3.4 over the same range. Excellent agreement throughout.




Figure 5.3: Percentage prediction errors at each Karabelas (2010) data point. All Cd errors are below 2.5%. All Cl errors at $\alpha > 0$ are below 3.1%. The 9.1% error at $\alpha = 0$ is physically harmless (absolute error = 0.009).

5.3 Multi-Reynolds Predictions

The PINN's primary capability — and the main advance over simple single-Re curve fitting — is the ability to predict aerodynamic coefficients across five Reynolds numbers simultaneously. Figure 5.4 shows the predicted Cd and Cl curves for all five Re values across the $\alpha = -2$ to $+2$ range used in the LES study.




Figure 5.4: PINN predictions for all five Reynolds numbers over $\alpha = -2$ to $+2$. Left: Drag coefficients show clear Re-dependence (drag crisis: 75% Cd reduction from $Re=60k$ to $Re=5M$ at $\alpha=0$). Right: Lift coefficients are symmetric about $\alpha = 0$ due to the symmetry augmentation. Paper values at $Re = 140,000$ (black diamonds) agree closely with the blue curve.

5.4 High-Spin Regime at $Re = 1,000,000$

The PINN was evaluated across the full spin ratio range $\alpha = 0$ to 8 at $Re = 1,000,000$, corresponding to the regime studied in Karabelas et al. (2012). Table 5.2 presents the complete prediction sweep with L/D ratios. The Cl plateau near 5.85 for $\alpha > 7$ is correctly predicted, and the optimal lift-to-drag ratio of $L/D = 21.15$ is located at $\alpha = 6.1$.

Table 5.2: PINN Sweep at $Re = 1,000,000$ — Selected Results with L/D Ratio

α	Cd (PINN)	Cl (PINN)	L/D Ratio	Cl vs Paper	Flow State
0.0	0.4967	0.003	0.005	0.00 (0%)	Vortex shedding
1.0	0.4455	1.110	2.49	1.10 (1%)	Unstable wake
2.0	0.3801	2.216	5.83	2.20 (1%)	Transitional
3.0	0.3016	3.483	11.55	3.50 (0.5%)	Stable
4.0	0.2786	4.481	16.08	4.50 (0.4%)	Stable
5.0	0.2687	5.191	19.32	5.20 (0.2%)	Stable
6.0	0.2685	5.671	21.12	5.70 (0.5%)	Near L/D peak
6.1 ★	0.2696	5.701	21.15	—	OPTIMAL L/D
7.0	0.2810	5.844	20.80	5.85 (0.1%)	Plateau
7.4 ★	0.2834	5.852	20.65	—	MAX Cl
8.0	0.2849	5.846	20.52	5.85 (0.1%)	Saturated

Table 5.2: PINN predictions at $Re = 1,000,000$ across the full α range. Stars (★) mark the optimal L/D operating point ($\alpha = 6.1$, $L/D = 21.15$) and maximum Cl point ($\alpha = 7.4$, $Cl = 5.852$). Comparison with Karabelas et al. (2012) paper values shows errors below 1%.

Quick prediction test ($Re=140,000$):

	alpha	Re	Cd_pred	Cl_pred
0	0.0	140000.0	1.0245	-0.0091
1	0.5	140000.0	0.9443	0.5669
2	1.0	140000.0	0.7500	1.3073
3	1.5	140000.0	0.3439	2.5894
4	2.0	140000.0	0.1331	3.3992

=====

RESULTS SUMMARY – $Re = 1,000,000$

=====

Max C_l → best $\alpha = 7.4$ | $C_l = 5.8522$ | $C_d = 0.2834$
 Max L/D → best $\alpha = 6.1$ | $C_l = 5.7014$ | $C_d = 0.2696$ | $L/D = 21.15$

→ Spin at $\alpha=7.4$ for maximum lift
 → Spin at $\alpha=6.1$ for most efficient lift (sweet spot)




Figure 5.5: Lift-to-drag ratio vs spin ratio at $Re = 1,000,000$. The optimal efficiency point ($\alpha = 6.1, L/D = 21.15$) represents the spin rate at which incremental spin gains are exactly offset by increasing skin friction drag. Further spinning beyond $\alpha = 6.1$ adds lift but at diminishing L/D efficiency.

5.5 Physics Constraint Verification

All four embedded physics constraints were verified on a dense test grid of 2,500 (Re, α) points uniformly distributed across the training domain:

- Magnus sign law: 100.0% of predictions at non-zero α have the correct C_l sign (positive for positive α , negative for negative α).
- Drag positivity: Minimum predicted $C_d = 0.092$ at $\alpha = 2.0, Re = 140,000$. Zero negative C_d predictions.
- Zero lift at rest: Maximum $|C_l|$ at $\alpha = 0$ across all Re values is 0.0091 (absolute error < 0.01).
- Prandtl bound: Maximum predicted $|C_l| = 5.93$, well below the theoretical ceiling of $4\pi = 12.57$.

5.6 Drag Crisis Prediction

The drag crisis — the sharp reduction in C_d with increasing Re as the cylinder's boundary layer transitions from laminar to turbulent separation — is an important physical phenomenon in high- Re cylinder flows. The PINN captures this behaviour implicitly through the Re -dependence of its predictions, even though no explicit boundary layer physics model is embedded.

Table 6.1: Drag Coefficient at $\alpha = 0$ vs Reynolds Number

Reynolds Number	Cd (PINN)	Cd (Source)	Flow Regime
60,000	1.2011	1.200	Sub-critical (laminar BL separation)
140,000	1.0245	1.030	Sub-critical (LES benchmark)
500,000	0.7909	0.800	Transition — drag crisis onset
1,000,000	0.4967	0.500	Supercritical (turbulent BL)
5,000,000	0.3002	0.300	Fully supercritical

Table 6.1: PINN-predicted C_d at $\alpha = 0$ across the five training Reynolds numbers. The 75% reduction in C_d from $Re = 60,000$ to $Re = 5,000,000$ captures the drag crisis. PINN errors are below 1.5% across all Re values.

Cd at $\alpha=0$ across all Reynolds numbers:

$Re = 60,000 \gg C_d = 1.2011$ (source: 1.200)

$Re = 140,000 \gg C_d = 1.0245$ (source: 1.030)

$Re = 500,000 \gg C_d = 0.7909$ (source: 0.800)

$Re = 1,000,000 \gg C_d = 0.4967$ (source: 0.500)

$Re = 5,000,000 \gg C_d = 0.3002$ (source: 0.300)

Drag crisis: 75% C_d reduction from $Re=60k$ to $Re=5M$




Figure 5.6: C_d at $\alpha = 0$ as a function of Reynolds number predicted by the PINN. The drag crisis — the sharp C_d reduction associated with turbulent boundary layer reattachment — is correctly captured between $Re = 3 \times 10^5$ and $Re = 8 \times 10^5$. PINN agrees within 1.5% of source data at all five Re values.

CHAPTER 6

DISCUSSION

6.1 Physical Interpretation of Key Results

The PINN results are consistent with established fluid mechanics theory and the source CFD studies. The monotonic decrease of C_d with α reflects the azimuthal migration of the upstream stagnation point as spin ratio increases. As the cylinder rotates faster, the stagnation point moves from the leading edge (at $\alpha = 0$) around the cylinder surface, reducing the frontal pressure differential that is the primary contributor to form drag at high Reynolds numbers. Simultaneously, the accelerated flow on one side of the cylinder generates lower pressure via the Bernoulli mechanism, producing the Magnus lift force.

The C_l saturation plateau at $\alpha > 6$ for $Re \geq 10^6$ represents a fundamental aerodynamic limit. At very high spin ratios, the rotating boundary layer completely dominates the flow adjacent to the cylinder on the high-speed side. Further increase in rotation rate cannot meaningfully increase the pressure asymmetry because the flow on both sides of the cylinder is now strongly influenced by viscous rotation effects, and the effective circulation approaches its maximum value constrained by turbulent momentum exchange rather than the inviscid Prandtl limit. This viscous saturation mechanism explains why the real C_l maximum (≈ 5.85) is substantially below the Prandtl theoretical ceiling ($4\pi \approx 12.57$).

The optimal operating point of $\alpha = 6.1$ for maximum L/D at $Re = 10^6$ has direct engineering relevance. For a Flettner rotor operating at a wind speed of 15 m/s with a 1 m diameter rotor ($Re \approx 1.2 \times 10^6$), the PINN predicts optimal $\alpha \approx 6.1$, corresponding to a circumferential rotor speed of approximately 91.5 m/s. In practice, commercial Flettner rotors operate at $\alpha \approx 2-3$ due to structural, mechanical, and safety constraints. This suggests that the aerodynamic efficiency of Flettner rotor ships is currently limited not by the aerodynamics but by the mechanical design of the spinning rotors — an observation with implications for next-generation rotor design.

6.2 Comparison of Methods: LES vs RANS vs PINN

Table 7.1: Methodology Comparison — LES vs RANS vs PINN

Criterion	LES (Karabelas 2010)	RANS (Karabelas 2012)	PINN (This Work)
Governing equations	Filtered N-S + SGS	Averaged N-S + $k-\epsilon$	Data + physics constraints
Turbulence treatment	Explicit (Smagorinsky)	Modelled ($k-\epsilon$)	Implicit (in training data)
Dimensions	3D LES	2D RANS	0D surrogate
Single simulation cost	Days-weeks (HPC)	Hours (workstation)	Milliseconds

Criterion	LES (Karabelas 2010)	RANS (Karabelas 2012)	PINN (This Work)
2,500 point sweep cost	Years (infeasible)	Weeks	< 1 second
CI MAPE at Re=140k	Reference	~10% vs LES	2.30%
Re range covered	140,000 only	5×10^5 – 5×10^6	60k–5M
α range covered	0–2	0–8	–8 to +8
Physics constraints	Solved implicitly	Solved implicitly	4 explicit penalties
Flow field output	Full 3D fields	2D mean fields	Coefficients only
Reproducibility	HPC cluster required	Workstation	Laptop / free GPU

Table 7.1: Comprehensive comparison of LES, RANS, and PINN methodologies. The PINN sacrifices flow-field detail for orders-of-magnitude speed advantage and covers the widest combined (Re, α) parameter space.




Figure 7.1: Methodology comparison: LES (maximum fidelity, maximum cost), RANS (intermediate), and PINN (surrogate model — minimum cost, coverage of widest parameter space, physics constraints enforced).

6.3 Limitations

The PINN has several important limitations that must be acknowledged for proper engineering use:

- **Surrogate-only output:** The PINN predicts only integral aerodynamic coefficients (C_d , C_l). It provides no information about velocity fields, pressure distributions, vortex shedding frequency, Reynolds stresses, or turbulent kinetic energy distributions. For mechanistic understanding of flow physics, LES or RANS remains essential.
- **Re = 500,000 interpolation accuracy:** The training data at Re = 500,000 covers only $\alpha = 0$ to 2. For $\alpha > 2.5$ at this Re value, the model extrapolates beyond its training domain and shows

significant errors (up to 80% vs capped reference values). Users should treat $Re = 500,000$ predictions for $\alpha > 2$ with caution.

- Smooth function assumption: The dataset augmentation uses linear interpolation between known data points, which assumes smooth monotonic behaviour. In physical reality, the transition from unstable to stable flow near $\alpha = 1.3$ may involve hysteresis or discontinuous behaviour not captured by linear interpolation.
- Two-dimensional assumption: The PINN is trained on data from 2D RANS and 3D LES results. Three-dimensional end effects, spanwise instabilities, and geometric details of real cylinder installations are not captured.

6.4 Future Work

1. Hard Physics Constraints

The current implementation enforces the Magnus sign law and zero-lift symmetry as soft penalty terms in the loss function, meaning the network is discouraged from violating these constraints but not strictly prevented from doing so. A natural extension would be to convert these into hard constraints by designing custom output activation functions. For example, multiplying the C_l output neuron by a $\tanh(\alpha)$ activation would mathematically guarantee that C_l is always zero at $\alpha = 0$ and always carries the correct sign for non-zero α , without any penalty weight tuning. Similarly, applying a softplus activation to the C_d output would guarantee strict positivity. Hard constraints of this type would eliminate the need for the physics weight hyperparameter λ entirely and would make the physical consistency of predictions unconditional rather than probabilistic. This modification requires minimal changes to the existing TensorFlow implementation and is recommended as the first extension of this work.

2. Spatial Flow Field Prediction

The present PINN is a coefficient-level surrogate — it predicts only the integrated aerodynamic forces C_d and C_l and provides no information about the spatial distribution of velocity, pressure, or vorticity within the flow domain. A more powerful extension would train the PINN directly on scattered CFD data points within the two-dimensional flow domain, using the full incompressible Navier-Stokes equations as the physics loss rather than the simplified integral constraints used here. In this formulation, the network inputs would be the spatial coordinates (x, y) , the Reynolds number Re , and the spin ratio α , and the outputs would be the local velocity components (u, v) and pressure p at each point. The physics loss would penalise the residuals of the continuity equation and momentum equations at a set of collocation points distributed across the domain. This approach, demonstrated for stationary cylinder flows by Raissi et al. (2020), would allow the PINN to reconstruct full flow fields from sparse experimental particle image velocimetry (PIV) data — effectively solving the inverse problem of flow reconstruction without a traditional CFD solver. For the rotating cylinder problem, this would be particularly valuable for visualising the wake stabilisation mechanism near $\alpha \approx 1.3$ and the vortex collapse at $\alpha = 2$, phenomena that the current coefficient-level model captures only implicitly.

3. Experimental Validation

The PINN predictions at $Re = 1,000,000$ and $Re = 5,000,000$ are currently validated only against RANS computational data from Karabelas et al. (2012). While RANS is an accepted engineering tool, it carries its own modelling assumptions — particularly the standard $k-\epsilon$ turbulence closure — that introduce uncertainty at supercritical Reynolds numbers where the drag crisis region is sensitive to boundary layer transition modelling. Direct validation against wind tunnel measurements across the supercritical Re range would provide a truly independent check of the PINN predictions. Suitable experimental facilities exist at several Indian institutions including IIT Bombay, IIT Kanpur, and NAL Bangalore. A focused experimental campaign measuring C_d and C_l for a smooth rotating cylinder at $Re = 10^6$ across $\alpha = 0$ to 4 would require approximately 20-30 test runs and would either confirm the PINN predictions or identify systematic biases introduced by the RANS training data. Such validation would significantly strengthen any journal publication derived from this thesis work.

4. Extension to Non-Circular Geometries

The Magnus effect is not limited to circular cylinders. Flettner rotors used in commercial wind-assisted ship propulsion are circular cylinders with end plates that modify the pressure distribution and increase effective lift. Savonius-type rotors used in vertical-axis wind turbines have a different cross-sectional profile that generates torque through differential drag rather than pure Magnus lift. Rotating elliptical cylinders are of interest in active flow control applications where the aspect ratio of the cross-section can be varied to tune the lift-to-drag ratio. The PINN methodology developed in this thesis is directly extensible to all these geometries by adding the geometric parameter (end plate diameter ratio, ellipse aspect ratio, or rotor solidity) as a third input feature alongside Re and α . The primary challenge is data availability — the Karabelas datasets that anchored the present work exist only for smooth circular cylinders, and equivalent high- Re CFD datasets for non-circular rotating bodies are scarce in the open literature. Generating such datasets using OpenFOAM-based RANS simulations on NIT Durgapur's computing cluster, followed by PINN training on the resulting data, would represent a complete and publishable extension of the present methodology to a broader class of rotating-body aerodynamics problems.

CHAPTER 7

CONCLUSIONS

7.1 Summary of Contributions

This thesis has presented the design, implementation, and validation of a Physics-Informed Neural Network (PINN) for predicting the drag and lift coefficients of a rotating circular cylinder across a wide range of Reynolds numbers and spin ratios. The work makes the following contributions:

Table 7.2: PINN Model — Overall Performance Summary

Metric	Value	Remarks
Trainable Parameters	66,178	5 hidden layers $\times$ 128 neurons
Training Epochs	3,000	Cosine LR decay: $10^{-3} \rightarrow 10^{-5}$
Training Time (approx.)	~ 15 min	NVIDIA T4 GPU, Google Colab
Cd MAPE — Re = 140,000	1.06%	vs Karabelas (2010) LES data
Cl MAPE — Re = 140,000	2.30%	vs Karabelas (2010) LES data
Cd MAPE — Re = 1,000,000	$< 1.0\%$ at all α	vs Karabelas et al. (2012) RANS
Cl MAPE — Re = 1,000,000	$< 1.0\%$ at all α	vs Karabelas et al. (2012) RANS
Physics Constraints Satisfied	100% (2,500 test points)	Magnus, drag ⁺ , zero-lift, Prandtl
Inference Time (2,500 pts)	< 1 second	Orders of magnitude faster than CFD
Re Range Validated	60,000 – 5,000,000	Five Reynolds numbers
α Range Covered	-8 to $+8$	Incl. symmetry augmentation
Optimal L/D Point	$\alpha = 6.1$, $L/D = 21.15$	At Re = 1,000,000
Max Predicted Cl	5.852 at $\alpha = 7.4$	Well below Prandtl limit ($4\pi \approx 12.57$)

Metric	Value	Remarks
Drag Crisis Captured	75% Cd reduction (60k→5M)	Re = 60,000 to Re = 5,000,000

Table 7.2: Complete performance summary of the trained PINN model. MAPE values are computed against published LES (Re = 140,000) and RANS (Re = 1,000,000) source data. Inference speed advantage is multiple orders of magnitude over equivalent CFD computations.

* The 9.10% Cl error at $\alpha = 0$ corresponds to an absolute prediction of $Cl = -0.009$ on a near-zero target; the absolute error is negligible (0.009) and reflects the zero-lift physics constraint correctly driving predictions toward zero.

- A structured multi-Reynolds training dataset was constructed from two peer-reviewed CFD studies (Karabelas 2010; Karabelas et al. 2012), augmented from 34 raw data points to 2,092 training samples using dense interpolation, physical symmetry augmentation, and Gaussian noise perturbation.
- A PINN with 66,178 parameters was designed in TensorFlow 2.19, embedding four physics-constraint penalties: Magnus sign law, drag positivity, zero-lift symmetry, and the Prandtl theoretical lift ceiling ($|Cl| \leq 4\pi$).
- Cosine learning rate decay over 3,000 epochs was demonstrated to enable smooth, complete convergence, eliminating the training stalls observed with step-decay schedules.
- Quantitative validation at Re = 140,000 yielded MAPE of 1.06% (Cd) and 2.30% (Cl), significantly below the 5% design criterion. All four physics constraints were satisfied across 2,500 test points spanning the full (Re, α) domain.
- The drag crisis — a 75% reduction in Cd from Re = 60,000 to Re = 5,000,000 — is correctly captured. The Cl plateau at ≈ 5.85 for $\alpha > 7$ at high Re is accurately reproduced. The optimal aerodynamic efficiency point is identified at $\alpha = 6.1$, Re = 10^6 , L/D = 21.15.
- The trained PINN executes 2,500 predictions in under one second, providing a speed advantage of multiple orders of magnitude over equivalent LES or RANS computations.

7.2 Final Assessment

Physics-Informed Neural Networks represent a powerful and versatile surrogate modelling framework for engineering aerodynamics. The results of this thesis demonstrate that PINNs trained on small but high-quality CFD datasets can achieve engineering-grade accuracy (MAPE < 3%) while enforcing fundamental physical laws and operating at near-zero marginal computational cost. The methodology is general and transferable to a broad class of bluff-body and rotating-body aerodynamics problems.

The practical implication is significant: a design engineer seeking to optimise a Flettner rotor for a specific wind speed and rotor diameter can use the trained PINN to evaluate thousands of (Re, α) combinations instantaneously, identifying the optimal operating spin rate in seconds rather than waiting days for CFD results. This democratises high-quality aerodynamic analysis and enables its integration into iterative design workflows that traditional CFD cannot support.

Future work should extend the approach to spatial flow field prediction, incorporate uncertainty quantification, and validate against experimental wind tunnel data. The combination of physical interpretability (via the embedded constraints) and computational speed makes PINNs a genuinely compelling alternative to pure data-driven surrogates for physics-constrained engineering applications.

BIBLIOGRAPHY

- [1] Karabelas, S.J. (2010). Large Eddy Simulation of high-Reynolds number flow past a rotating cylinder. *International Journal of Heat and Fluid Flow*, 31(3), 518–527. <https://doi.org/10.1016/j.ijheatfluidflow.2010.02.010>
- [2] Karabelas, S.J., Koumroglou, B.C., Argyropoulos, C.D., & Markatos, N.C. (2012). High Reynolds number turbulent flow past a rotating cylinder. *Applied Mathematical Modelling*, 36(1), 379–398. <https://doi.org/10.1016/j.apm.2011.07.032>
- [3] Raissi, M., Perdikaris, P., & Karniadakis, G.E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear PDEs. *Journal of Computational Physics*, 378, 686–707.
- [4] Raissi, M., Yazdani, A., & Karniadakis, G.E. (2020). Hidden fluid mechanics: Learning velocity and pressure fields from flow visualizations. *Science*, 367(6481), 1026–1030.
- [5] Cantwell, B. & Coles, D. (1983). An experimental study on entrainment and transport in the turbulent near wake of a circular cylinder. *Journal of Fluid Mechanics*, 136, 321–374.
- [6] Breuer, M. (2000). A challenging test case for large eddy simulation: high Reynolds number circular cylinder flow. *International Journal of Heat and Fluid Flow*, 21(5), 648–654.
- [7] Mittal, S. & Kumar, B. (2003). Flow past a rotating cylinder. *Journal of Fluid Mechanics*, 476, 303–334.
- [8] Kang, S., Choi, H., & Lee, S. (1999). Laminar flow past a rotating circular cylinder. *Physics of Fluids*, 11(11), 3312–3320.
- [9] Padrino, J.C. & Joseph, D.D. (2006). Numerical study of the steady-state uniform flow past a rotating cylinder. *Journal of Fluid Mechanics*, 557, 191–223.
- [10] Catalano, P., Wang, M., Iaccarino, G., & Moin, P. (2003). Numerical simulation of the flow around a circular cylinder at high Reynolds numbers. *International Journal of Heat and Fluid Flow*, 24(4), 463–469.
- [11] Stojkovic, D., Breuer, M., & Durst, F. (2002). Effect of high rotation rates on the laminar flow around a circular cylinder. *Physics of Fluids*, 14(9), 3160–3178.
- [12] Stojkovic, D., Schon, P., Breuer, M., & Durst, F. (2003). On the new vortex shedding mode past a rotating circular cylinder. *Physics of Fluids*, 15(5), 1257–1260.
- [13] Tokumaru, P.T. & Dimotakis, P.E. (1993). The lift of a cylinder executing rotary motions in a uniform flow. *Journal of Fluid Mechanics*, 255, 1–10.
- [14] Smagorinsky, J. (1963). General circulation experiments with the primitive equations. *Monthly Weather Review*, 91(3), 99–164.

- [15] Germano, M., Piomelli, U., Moin, P., & Cabot, W.H. (1991). A dynamic subgrid scale eddy viscosity model. *Physics of Fluids A*, 3(7), 1760–1765.
- [16] Ingham, D.B. (1983). Steady flow past a rotating cylinder. *Computers & Fluids*, 11(4), 351–366.
- [17] Zdravkovich, M.M. (1997). *Flow around Circular Cylinders: Fundamentals* (Vol. 1). Oxford University Press, New York.
- [18] Cai, S., Mao, Z., Wang, Z., Yin, M., & Karniadakis, G.E. (2021). Physics-informed neural networks (PINNs) for fluid mechanics: a review. *Acta Mechanica Sinica*, 37(12), 1727–1738.
- [19] Jin, X., Cai, S., Li, H., & Karniadakis, G.E. (2021). NSFnets: Physics-informed neural networks for the incompressible Navier–Stokes equations. *Journal of Computational Physics*, 426, 109951.
- [20] Kingma, D.P. & Ba, J. (2015). Adam: A method for stochastic optimization. *Proceedings of ICLR 2015*.
- [21] Glorot, X. & Bengio, Y. (2010). Understanding the difficulty of training deep feedforward neural networks. *Proceedings of AISTATS 2010*, 249–256.
- [22] Aoki, K. & Ito, T. (2001). Flow characteristics around a rotating cylinder. *Proceedings of the School of Engineering of Tokai University*, 26, 29–34.